Discussion

Assume we want to learn the relationship between education and income, in a sample of private data.

- Sketch out three differentially private approaches to learning this? (Make any assumptions you need but write them down.)
- Which of your methods would work if we further extended the relationship to many features/variables/covariates?
- **If time:** Does it matter if our model is predictive or inferential?
- **If time:** What would be an attack on this model if it were released without privacy-preservation?